

FunTAL: Reasonably Mixing a Functional Language with Assembly

Daniel Patterson

Northeastern University
dbp@ccs.neu.edu

Jamie Perconti

Northeastern University
jamestperconti@gmail.com

Christos Dimoulas

Harvard University
chrdimo@seas.harvard.edu

Amal Ahmed

Northeastern University
amal@ccs.neu.edu

Abstract

We present FunTAL, the first multi-language system to formalize safe interoperability between a high-level functional language and low-level assembly code while supporting compositional reasoning about the mix. A central challenge in developing such a multi-language is how to bridge the gap between assembly, which is staged into calls to continuations, and high-level code, where subterms return a result. We present a *compositional* stack-based typed assembly language that supports *components*, comprised of one or more basic blocks, that may be embedded in high-level contexts. We also present a logical relation for FunTAL that supports reasoning about equivalence of high-level components and their assembly replacements, mixed-language programs with callbacks between languages, and assembly components comprised of different numbers of basic blocks.

This work is a first step towards multi-language formalisms that (1) allow developers to replace high-level components with performant low-level implementations while reasoning about safety and correctness of the replacement; (2) can be used to specify compiler correctness theorems that permit compiled components to be linked with low-level code compiled from other languages (as proposed by Perconti and Ahmed [? ?]); and (3) can be used to verify just-in-time compilers which replace portions of high-level code with equivalent assembly, inserting callbacks back to high-level code.

Note: We use **green** to typeset our functional language \mathbb{F} and **purple** to typeset our typed assembly language \mathbb{T} . This paper will be difficult to follow unless read/printed in color.

1. Introduction

Developers frequently integrate code written in lower-level languages into their high-level-language programs. For instance, OCaml and Haskell developers may leverage the FFI to make use of libraries implemented in C, while Rust de-

velopers can include inline assembly. In each of these cases, developers resort to the lower-level language so they can use features unavailable in the high-level language and gain access to hardware or fine-tune performance.

However, the benefits of mixed-language programs come at a price. To reason about the behavior of a high-level component, developers need to think not only about the semantics of the high-level language, but also about all potential interactions between high- and low-level code and the way their high-level code was compiled. Since low-level code comes without safety guarantees, invalid instructions could crash the program. More insidiously, low-level code can potentially alter control flow, mutate values that it shouldn't have access to, or introduce security vulnerabilities that wouldn't be possible in the source language. Unfortunately, there are no mixed-language systems that can aid programmers who occasionally wish to work at a lower level of abstraction—i.e., systems that allow programmers to specify lower-level invariants (when they are aware of them), guarantee safe interoperability, and provide rules for compositional reasoning in a mixed-language setting.

Even if developers don't directly write inline assembly, mixed-language programs are a reality that compiler writers and compiler-verification efforts must contend with. For instance, mixed programs show up in modern just-in-time (JIT) compilers, where the high-level language is initially interpreted until the runtime can identify portions to statically compile, at which point those portions of the code are replaced with equivalent assembly. These assembly components will include hooks to move back into the interpreted runtime, corresponding closely to the semantics of a multi-language program. Verifying correctness of such JITs requires proving that the high-level fragment and its compiled version are equivalent in the mixed language.

In the case of traditional compilers, compiled components are frequently linked with target code compiled from a dif-

ferent source language, or with low-level routines that form part of the runtime system. Perconti and Ahmed [?] argue that correctness theorems for verified compilers that account for such linking require reasoning about mixed-language programs. Specifically, they set up a multi-language that specifies the rules of source-target interoperability and then express compiler correctness as a multi-language equivalence between a source component e_S and its compiled version e_T . Hence, the theorem ensures that e_T linked with some arbitrary target code e'_T will behave the same as e_S interoperating with e'_T .

The above scenarios call for the design of a multi-language that specifies interoperability between a high-level language and assembly, along with proof methods for reasoning about equivalence of components in this multi-language. Note that Perconti and Ahmed [?] left the design of a multi-language that embeds assembly as future work, to be tackled when they verify their code generation pass. Their functional-language compiler performs only closure conversion and heap allocation so their multi-language need only embed languages that naturally support compositional reasoning.

We investigate the design of a multi-language system that allows assembly to be embedded in a typed functional language and vice versa. A key difficulty is ensuring that the embedded assembly has local and well-controlled effects. This is challenging because assembly is inherently *non-compositional*—control can change to an arbitrary point with direct jumps and code can access arbitrary values far up on the call-stack. To allow a compositional functional language to safely interoperate with assembly, such behavior must be constrained, which we do using types at the assembly level. Moreover, we need to identify the right notion of *component* in assembly: intuitively, an assembly component may be comprised of multiple basic blocks and we should be able to show equivalence between terms of the functional language (i.e., high-level components) and multi-block assembly components. But how do we identify which blocks should be grouped together into a component without imposing so much high-level structure on assembly that it ceases to be low level? Even once we identify such groupings, we must still contend with the control-flow gap between a direct-style functional language in which terms return results and assembly code that is staged into jumps to continuations. Finally, we must find a way to embed functional code in assembly so we can support callbacks from assembly to the functional language.

Contributions We make the following contributions:

- We design a compositional typed assembly language (TAL) called \mathbb{T} , building on the stack-based typed assembly language of Morrisett *et al.* [?] (henceforth, STAL). The central novelty of our TAL \mathbb{T} are extensions to an STAL-like type system that help us reason about multi-block components and bridge the gap between direct-

style high-level components and continuation-based assembly components (§3).

- We present a multi-language \mathbb{FT} in the style of Matthews-Findler [?] that supports interoperability between a simply typed functional language \mathbb{F} with recursive types and our TAL \mathbb{T} (§4).
- We develop a novel step-indexed Kripke logical relation for reasoning about equivalence of \mathbb{FT} components (§5). It builds on prior logical relations for mutable state [?], but is the first to support reasoning about equivalence of programs that mix assembly with lambdas (including callbacks between them), and of assembly components comprised of different numbers of basic blocks. The central novelty lies in the mechanics of accommodating assembly and equivalence of multi-block components.

The supplementary material includes complete language semantics, definitions, and proofs.

2. Main Ingredients of the Mix

We design a *compositional* TAL \mathbb{T} that draws largely from Morrisett *et al.*'s STAL [?], which has a single explicit stack and assembly instructions to allocate, read, write, and free stack cells. We follow much of their basic design, including using stack-tail polymorphism to hide values on the stack so they will be preserved across calls, and the use of register-file and stack typing to specify preconditions for jumping to a code block.

Our main novelty is identifying the notion of a TAL *component*. In \mathbb{T} , we need to be able to reason about a component e_T , because we will eventually be embedding these components as terms in a high-level functional language called \mathbb{F} . A component e_T must be composed of assembly instructions, but we don't want to restrict it to a single basic block, so we use a pair (\mathbf{I}, \mathbf{H}) of an instruction sequence \mathbf{I} and a local heap fragment \mathbf{H} that maps locations to code blocks used in local intra-component jumps.

The combined language \mathbb{FT} is a typical Matthews-Findler multi-language [?], where the syntax of both languages are combined and boundary terms are added to mediate interactions between the two. A boundary term $\tau \mathcal{FT} e_T$ means that the \mathbb{T} component e_T within the boundary will be used in an \mathbb{F} context at type τ . To be well-typed, the inner component e_T should have the type translated from τ according to the multi-language type translation in §4.

\mathbb{FT} exists to enable reasoning about the equivalence of \mathbb{F} expressions and \mathbb{T} components, or mixed combinations of the two. Intuitively, we would like to treat blocks of assembly as similar to functions in high-level languages. Semantically, functions are objects that, given related inputs, produce related outputs. Following STAL we can, at least, model the state of the stack and a subset of the registers as inputs. But blocks of assembly instructions do not have a

clear statically determinable output, leading us towards one of our central novel contributions.

In STAL, every basic block has type $\forall[\Delta].\{\chi; \sigma\}$, where Δ contains type parameters, and χ and σ are respectively the register and stack typing preconditions. Since every block is in continuation style, blocks never return, always jumping to the next block, so there never need be outputs to relate — the output of a block is just the input constraints on the block to which it jumps. In our mixed-language setting we must, therefore, provide components with return continuations which, when called from high-level code, contain a halting instruction, and when called from assembly, jump to the next step in execution. In order to determine result types, we extend the STAL code pointer type to $\forall[\Delta].\{\chi; \sigma\}^q$, where q is our critical addition.

A *return marker* q specifies the register or stack position where the return continuation is stored, which allows us, following a basic calling convention, to determine the type of the value that will be passed to that continuation. As we'll see in later sections, there are a few other forms that q can take, but they all support our ability to reason about \mathbb{T} components as semantic objects that produce values of a specific type. This allows us to reason not only about the equivalence of structurally different assembly components made up of different numbers of basic blocks, but of components made up of entirely different mixes of languages.

3. Typed Assembly Language: \mathbb{T}

Syntax Figure 1 presents the full syntax of \mathbb{T} , our typed assembly language. Value types τ are the types ascribed to values small enough to fit in a register, including base values, recursive and existential types, and mutable (**ref**) or immutable (**box**) pointers to heap values. Word values w include unit $()$, integers n , locations ℓ , existential **packs**, and recursive **fold**s. We additionally follow STAL's convention that a word value w applied to a type instantiation ω is itself a value $w[\omega]$. Small values u include word values w , but also can be a register r that contains a word value. Instructions accept small values u as operands; hence, in the operational semantics, if u is a register we first load the value from the register, while if u is a word value we use it directly.

We ascribe heap-value types ψ to values h stored in the heap. These include tuples of word values $\langle w, \dots, w \rangle$ and code blocks $\text{code}[\Delta]\{\chi; \sigma\}^q.I$, which have types $\langle \tau, \dots, \tau \rangle$ and $\forall[\Delta].\{\chi; \sigma\}^q$, respectively. Note we have mutable **ref** references to tuples but only immutable **box** references to code, since we prohibit self-modifying code.

Code blocks $\text{code}[\Delta]\{\chi; \sigma\}^q.I$ specify a type environment Δ , a register file typing χ , and a stack type σ for an instruction sequence I . Here χ and σ are preconditions for safely jumping to I : χ is a mapping from registers r to the type of values τ the register must contain, while σ is a list of value types on top of the stack that may end with an abstract stack-tail variable ζ . The type variables in Δ , which

Value type τ	$::=$	$\alpha \mid \text{unit} \mid \text{int} \mid \exists \alpha. \tau \mid \mu \alpha. \tau$ $\text{ref } \langle \tau, \dots, \tau \rangle \mid \text{box } \psi$
Word value w	$::=$	$() \mid n \mid \ell \mid \text{pack } \langle \tau, w \rangle \text{ as } \exists \alpha. \tau$ $\text{fold}_{\mu \alpha. \tau} w \mid w[\omega]$
Register r	$::=$	$r1 \mid r2 \mid \dots \mid r7 \mid ra$
Small value u	$::=$	$w \mid r \mid \text{pack } \langle \tau, u \rangle \text{ as } \exists \alpha. \tau$ $\text{fold}_{\mu \alpha. \tau} u \mid u[\omega]$
Type instantiation ω	$::=$	$\tau \mid \sigma \mid q$
Heap value type ψ	$::=$	$\forall[\Delta].\{\chi; \sigma\}^q \mid \langle \tau, \dots, \tau \rangle$
Heap value h	$::=$	$\text{code}[\Delta]\{\chi; \sigma\}^q.I \mid \langle w, \dots, w \rangle$
Register typing χ	$::=$	$\cdot \mid \chi, r : \tau$
Stack typing σ	$::=$	$\zeta \mid \bullet \mid \tau :: \sigma$
Return marker q	$::=$	$r \mid i \mid \epsilon \mid \text{end } \{\tau; \sigma\}$
Type env Δ	$::=$	$\cdot \mid \Delta, \alpha \mid \Delta, \zeta \mid \Delta, \epsilon$
Heap typing Ψ	$::=$	$\cdot \mid \Psi, \ell : {}^\nu \psi \text{ where } {}^\nu ::= \text{ref} \mid \text{box}$
Memory M	$::=$	(H, R, S)
Heap fragment H	$::=$	$\cdot \mid H, \ell \mapsto h$
Register file R	$::=$	$\cdot \mid R, r \mapsto w$
Stack S	$::=$	$\text{nil} \mid w :: S$
Instruction sequence I	$::=$	$\epsilon; I$ instruction sequencing $\text{jmp } u$ jump to u within same component $\text{call } u \{ \sigma, q \}$ jump to u , with return address at q $\text{ret } r \{ r_r \}$ jump back to code at r with result in r_r $\text{halt } \tau, \sigma \{ r_r \}$ halt with value type τ in register r_r
Single instruction ι	$::=$	$\text{aop } r_d, r_s, u$ store result of $\text{add} \mid \text{mul} \mid \text{sub}$ in r_d $\text{bnz } r, u$ jump to u if r contains 0 $\text{ld } r_d, r_s[i]$ load from i th position in tuple at r_s $\text{st } r_d[i], r_s$ store to i th position in mutable tuple at r_d $\text{ralloc } r_d, n$ alloc mutable n -tuple initialized from stack $\text{balloc } r_d, n$ alloc immutable n -tuple initialized from stack $\text{mv } r_d, u$ move value u into register r_d $\text{salloc } n$ allocate n stack cells with unit values $\text{sfree } n$ free n stack cells $\text{sld } r_d, i$ load i th stack value into r_d $\text{sst } i, r_s$ store r_s into i th stack slot $\text{unpack } \langle \alpha, r_d \rangle u$ unpack existential, binding to α, r_d $\text{unfold } r_d, u$ unfold recursive type
Component e	$::=$	(I, H)
Halt instruction v	$::=$	$\text{halt } \tau, \sigma \{ r_r \}$
Evaluation context E	$::=$	$([\cdot], \cdot)$

Figure 1. \mathbb{T} Syntax

may appear free in χ , σ , and \mathbf{I} , must be instantiated when we jump to the code block. If this code block is stored at location ℓ , and register \mathbf{r} contains ℓ , we can jump to it via $\text{jmp } \mathbf{r}[\bar{w}]$ where \bar{w} instantiates the variables in Δ . (We use vector notation, e.g., \bar{w} or $\bar{\tau}$, to denote a sequence.)

As discussed in §2, our code blocks include a novel return marker \mathbf{q} , which tells us where to find the current return continuation. Here \mathbf{q} can be a register \mathbf{r} in χ , or a stack index \mathbf{i} that is accessible in σ (i.e., the \mathbf{i} th stack slot is not hidden in the stack tail ζ). Return markers can also range over type variables ϵ which we use to abstract over return markers (as we explain below). There is also a special return marker $\text{end}\{\tau; \sigma\}$ which means that when the current component finishes it should halt with a value of type τ and stack of type σ . In \mathbb{T} , this would mean the end of the program with a `halt` instruction, but within a multi-language boundary, the same `halt` results in a transition to the high-level language.

A memory \mathbf{M} includes a heap \mathbf{H} which maps locations ℓ to heap values \mathbf{h} , a register file \mathbf{R} which maps registers \mathbf{r} to word values \mathbf{w} , and a stack \mathbf{S} which is a list of word values.

An instruction sequence \mathbf{I} is a list of instructions terminated by one of three jump instructions (`jmp`, `call`, `ret`) or the `halt` instruction. The distinction between jump instructions is a critical part of \mathbb{T} explored in depth later in this section. Individual instructions ι include many standard assembly instructions and are largely similar to STAL.

A component \mathbf{e} is a tuple (\mathbf{I}, \mathbf{H}) of instructions \mathbf{I} and a local heap fragment \mathbf{H} . The local heap fragment can contain multiple local blocks used by the component. We distinguish the `halt` instruction as a value \mathbf{v} , as it is the only \mathbb{T} instruction sequence that does not reduce.

Operational Semantics We specify a small-step operational semantics as a relation on memories \mathbf{M} and components \mathbf{e} : $\langle \mathbf{M} \mid \mathbf{e} \rangle \mapsto \langle \mathbf{M}' \mid \mathbf{e}' \rangle$. Operationally, we merge local heap fragments to the global heap and then use the evaluation context \mathbf{E} to reduce instructions according to relation: $\langle \mathbf{M} \mid \mathbf{I} \rangle \mapsto \langle \mathbf{M}' \mid \mathbf{I}' \rangle$. In \mathbb{T} , evaluation contexts \mathbf{E} are not particularly interesting, but when \mathbb{T} is embedded within the multi-language, \mathbf{E} will include boundaries. While Figure 1 includes operational descriptions of the instructions, the full semantics are standard and elided.

Type System In Figure 2 we present a selection of typing rules for \mathbb{T} . We elide various type judgments and well-formedness judgments for small values, heap fragments, register files, as they are standard, focusing instead on novel rules for instructions, instruction sequences, and components. Full details exist in our supplementary material.

An instruction ι is typed under a static heap Ψ , a type environment Δ , a register file typing χ , a stack typing σ , and return marker \mathbf{q} . The instruction may change any of these except the static heap. Critically, instructions are only well-typed under a non- ϵ return marker, as a block of instructions needs to know to where it is returning. Code pointers can have ϵ in their return marker, but by the time they are

$$\boxed{\Psi; \Delta; \chi; \sigma; \mathbf{q} \vdash \iota \Rightarrow \Delta'; \chi'; \sigma'; \mathbf{q}'} \quad \mathbf{q} \neq \epsilon$$

$$\frac{\Psi; \Delta; \chi \vdash \mathbf{u} : \tau \quad \mathbf{q} \neq \mathbf{r}_d}{\Psi; \Delta; \chi; \sigma; \mathbf{q} \vdash \text{mv } \mathbf{r}_d, \mathbf{u} \Rightarrow \Delta; \chi[\mathbf{r}_d : \tau]; \sigma; \mathbf{q}}$$

$$\frac{\chi(\mathbf{r}_s) = \tau}{\Psi; \Delta; \chi; \sigma; \mathbf{r}_s \vdash \text{mv } \mathbf{r}_d, \mathbf{r}_s \Rightarrow \Delta; \chi[\mathbf{r}_d : \tau]; \sigma; \mathbf{r}_d}$$

$$\boxed{\Psi; \Delta; \chi; \sigma; \mathbf{q} \vdash \mathbf{I}} \quad \mathbf{q} \neq \epsilon$$

$$\frac{\Psi; \Delta; \chi \vdash \mathbf{u} : \text{box } \forall \square. \{\chi'; \sigma'\}^q \quad \Delta \vdash \chi \leq \chi' \cdot [\Delta]; \chi; \sigma \vdash \mathbf{q}}{\Psi; \Delta; \chi; \sigma; \mathbf{q} \vdash \text{jmp } \mathbf{u}}$$

$$\frac{\chi(\mathbf{r}) = \text{box } \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \quad \chi(\mathbf{r}') = \tau}{\Psi; \Delta; \chi; \sigma; \mathbf{r} \vdash \text{ret } \mathbf{r} \{\mathbf{r}'\}}$$

$$\frac{\begin{array}{l} \Psi; \Delta; \chi \vdash \mathbf{u} : \text{box } \forall [\zeta, \epsilon]. \{\hat{\chi}; \hat{\sigma}\}^q \\ \Delta \vdash \hat{\chi} \setminus \hat{\mathbf{q}} \quad \text{ret-addr-type}(\hat{\mathbf{q}}, \hat{\chi}, \hat{\sigma}) = \forall \square. \{\mathbf{r} : \tau; \hat{\sigma}'\}^\epsilon \\ \Delta \vdash \tau \quad \Delta \vdash \hat{\sigma}'[\sigma_0/\zeta] \\ \Delta \vdash \forall \square. \{\hat{\chi}[\sigma_0/\zeta][i+k-j/\epsilon]; \hat{\sigma}[\sigma_0/\zeta][i+k-j/\epsilon]\}^q \\ \Delta \vdash \chi \leq \hat{\chi}[\sigma_0/\zeta][i+k-j/\epsilon] \\ \sigma = \tau_0 :: \dots :: \tau_j :: \sigma_0 \quad \hat{\sigma} = \tau_0 :: \dots :: \tau_j :: \zeta \\ j < i \quad \hat{\sigma}' = \tau'_0 :: \dots :: \tau'_k :: \zeta \end{array}}{\Psi; \Delta; \chi; \sigma; \mathbf{i} \vdash \text{call } \mathbf{u} \{\sigma_0, i+k-j\}}$$

$$\frac{\begin{array}{l} \Psi; \Delta; \chi \vdash \mathbf{u} : \text{box } \forall [\zeta, \epsilon]. \{\hat{\chi}; \hat{\sigma}\}^q \quad \Delta \vdash \hat{\chi} \setminus \hat{\mathbf{q}} \\ \text{ret-addr-type}(\hat{\mathbf{q}}, \hat{\chi}, \hat{\sigma}) = \text{box } \forall \square. \{\mathbf{r} : \tau; \hat{\sigma}'\}^\epsilon \\ \Delta \vdash \tau \quad \Delta \vdash \hat{\sigma}'[\sigma_0/\zeta] \\ \Delta \vdash \forall \square. \{\hat{\chi}[\sigma_0/\zeta][\text{end}\{\tau^*; \sigma^*\}/\epsilon]; \hat{\sigma}[\sigma_0/\zeta][\text{end}\{\tau^*; \sigma^*\}/\epsilon]\}^q \\ \Delta \vdash \chi \leq \hat{\chi}[\sigma_0/\zeta][\text{end}\{\tau^*; \sigma^*\}/\epsilon] \\ \sigma = \bar{\tau} :: \sigma_0 \quad \hat{\sigma} = \bar{\tau} :: \zeta \quad \hat{\sigma}' = \bar{\tau}' :: \zeta \end{array}}{\Psi; \Delta; \chi; \sigma; \text{end}\{\tau^*; \sigma^*\} \vdash \text{call } \mathbf{u} \{\sigma_0, \text{end}\{\tau^*; \sigma^*\}\}}$$

$$\boxed{\Psi; \Delta; \chi; \sigma; \mathbf{q} \vdash \mathbf{e} : \tau; \sigma'} \quad \mathbf{q} \neq \epsilon$$

$$\frac{\begin{array}{l} \Psi \vdash \mathbf{H} : \Psi' \quad \forall (\ell : \nu \psi) \in \Psi. \nu = \text{box} \\ \text{ret-type}(\mathbf{q}, \chi, \sigma) = \tau; \sigma' \quad (\Psi, \Psi'); \Delta; \chi; \sigma; \mathbf{q} \vdash \mathbf{I} \end{array}}{\Psi; \Delta; \chi; \sigma; \mathbf{q} \vdash (\mathbf{I}, \mathbf{H}) : \tau; \sigma'}$$

$$\begin{array}{l} \text{ret-type}(\mathbf{r}, \chi, \sigma) = \tau; \sigma' \text{ if } \chi(\mathbf{r}) = \text{box } \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \\ \text{ret-type}(\mathbf{i}, \chi, \sigma) = \tau; \sigma' \text{ if } \sigma(\mathbf{i}) = \text{box } \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \\ \text{ret-type}(\text{end}\{\tau; \sigma'\}, \chi, \sigma) = \tau; \sigma' \\ \text{ret-addr-type}(\mathbf{r}, \chi, \sigma) = \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \\ \text{if } \chi(\mathbf{r}) = \text{box } \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \\ \text{ret-addr-type}(\mathbf{i}, \chi, \sigma) = \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \\ \text{if } \sigma(\mathbf{i}) = \text{box } \forall \square. \{\mathbf{r}' : \tau; \sigma'\}^q \end{array}$$

Figure 2. Selected \mathbb{T} Typing Rules

jumped to, this must be instantiated. An example of this is shown later in this section.

The `mv` instruction shown in Figure 2 has two cases. In the first case, we are loading a small value u with type τ into register r_d , which we know is not the return marker q . The result of this is that the register file typing now reflects the updated register, and no other changes have occurred. The other case is that we are moving the value in register r_s into register r_d , and the former is the current return marker, so it is pointing to the return continuation. In that case, not only do we update the register file, we also change the return marker to reflect that the continuation is now in r_d . Other instructions, like `sst` and `sld`, similarly have cases depending on whether the operation will move the address of the return continuation.

Instruction sequences I are typed under the same Ψ , Δ , χ , σ , and q . In Figure 2, we show the three types of jump instructions. First, we can see an *intra-component jump*: the `jmp` instruction. This requires that the location u being jumped to be a code pointer that has preconditions χ' and σ for the register file and stack respectively. The current register file χ must be a subset of χ' , which means that we can have more registers with values in them, but the registers that occur in χ' must have matching types.

We also, critically, require that the return marker q on the code block being jumped to is the same as the current return marker. This captures the intuition of an intra-component jump. Finally, we have a restriction of form $\Delta'[\Delta]; \chi; \sigma \vdash q$ on the return marker. This judgment means that while type variables in Δ may be free in χ or σ , q must only include type variables in Δ' , which in this case is empty. This formally captures the notion that before a block can be jumped to, its return marker must have been fully instantiated, or informally that a block cannot abstract over its own return marker. A component can, however, have local blocks with abstract return markers. Consider the code pointer type:

$$\text{box } \forall[\epsilon]. \{ra: \text{box } \forall[] . \{r1: \tau; \sigma\}^\epsilon; \sigma\}^{ra}$$

This type is a pointer to a code block with a return marker type parameter ϵ that requires a stack of type σ and for register ra to be a code pointer. This inner code pointer is the continuation, as the entire block has ra as its return marker, but the return marker for this continuation is ϵ . When the continuation in ra is jumped to it requires that the stack still have type σ and that a value of type τ be stored in register $r1$. Since code pointers can't be jumped to until all their type variables are instantiated, the caller of this whole code block must provide a concrete continuation in register ra and instantiate ϵ with the corresponding concrete return marker before jumping.

The next instruction in Figure 2 is `ret`, which is the *inter-component* jump for returning from a component. Notably, the location being jumped to must be in a register; if it were still on the stack the type of σ would include itself. We require, first, that the register r being jumped to points to

a code block with no type variables, and second that the register r map to type τ , as required by the block being returned to. This is a type-enforced calling convention for the return value. Importantly, we make no restriction on the return marker q' on the block being jumped to. This is because with `ret` we are jumping back to a different component, which will in turn have its own return marker.

The last instruction shown in Figure 2 is `call`, which is our other *inter-component* jump. We include both typing rules for `call`, but focus on the case that the return marker is stack position i , as the other, with the halt marker $\text{end}\{\tau; \sigma\}$, is the same rule with fewer constraints.

Calls must protect their return continuation. In some assembly languages, there is a convention that certain registers (“callee-saved”) will be preserved such that when a call returns, they have the same value as before. However, we follow STAL in protecting values solely through stack-tail polymorphism, where a value can be stored in a part of the stack that has been abstracted away as a type variable. Static typing ensures that a callee that tried to read, write, or free values within the abstract tail would not type-check. Values that are accessible can be passed in front of the abstract tail, and the callee is free to allocate values in front, but typing constraints may force them to free the values before returning.

We can see that the index i where the return continuation is stored must be greater than j , the number of entries on the input stack σ in front of the tail σ_0 specified in the instruction. The location being jumped to, u , must be a code pointer with input registers $\hat{\chi}$ and stack $\hat{\sigma}$. Note that the prefix of $\hat{\sigma}$ matches the prefix of $\sigma, \tau_0 :: \dots :: \tau_j$, but $\hat{\sigma}$ has the abstract tail ζ .

The final formal parameter to call, $i + k - j$, is the return marker that the continuation for u must use. In particular, this is computed by taking the starting stack position, i , and then noting how the stack is modified between the input stack $\hat{\sigma}$ and output stack $\hat{\sigma}'$ by the code block pointed to by u . After the call, the stack has k values in front, but we know that position i was beyond the exposed j values, so the value on the stack at position i is now at position $i + k - j$.

The constraint that $\text{ret-addr-type}(\hat{q}, \hat{\chi}, \hat{\sigma})$ is $\forall[] . \{r: \tau; \hat{\sigma}'\}^\epsilon$ ensures that the block being jumped to has a return continuation where a value of type τ is stored in some register, the stack has type $\hat{\sigma}'$, and the return marker is ϵ . Operationally, u will get instantiated with $i + k - j$ for ϵ , which, based on the form of $\hat{\sigma}'$, means that the return continuation has preserved the original return location.

The register file subtyping constraint

$$\Delta \vdash \chi \leq \hat{\chi}[\sigma_0/\zeta][i+k-j/\epsilon]$$

ensures that the current register type χ is a subtype of the target $\hat{\chi}$ once it has been concretely instantiated with the stack tail and return address.

We similarly check with

$$\Delta \vdash \forall[] . \{\hat{\chi}[\sigma_0/\zeta][i+k-j/\epsilon]; \hat{\sigma}[\sigma_0/\zeta][i+k-j/\epsilon]\}^{\hat{q}}$$

```

f = (mv ra, l1ret; call l1 {•, end{int; •}}, H)
H(l1) = code[ζ, ε]{ra: ∇[].{r1: int; ζ}^ε; ζ}^ra.
      salloc 1; sst 0, ra; mv ra, l2ret[ζ, ε];
      call l2 {∇[].{r1: int; ζ}^ε :: ζ, 0}
H(l1ret) = code[] {r1: int; •} end{int; •}.
          halt int, • {r1}
H(l2) = code[ζ, ε]{ra: ∇[].{r1: int; ζ}^ε; ζ}^ra.
      mv r1, 1; jmp l2aux[ζ, ε]
H(l2aux) = code[ζ, ε]{r1: int, ra: ∇[].{r1: int; ζ}^ε; ζ}^ra.
      mult r1, r1, 2; ret ra {r1}
H(l2ret) = code[ζ, ε]{r1: int; ∇[].{r1: int; ζ}^ε :: ζ}^0.
      sld ra, 0; sfree 1; ret ra {r1}

```

Figure 3. \mathbb{T} Example: Call to Call




Figure 4. \mathbb{T} Control Flow: Call to Call (Fig. 3)

that the code block type is well-formed when concretely instantiated, and with $\Delta \vdash \hat{\sigma}'[\sigma_0/\zeta]$ that the resulting stack is well-formed once concretely instantiated. Finally, we ensure with $\Delta \vdash \hat{\chi} \setminus \hat{q}$ that if \hat{q} is a register, then $\hat{\chi}$ is well-formed without it. This means that while \hat{q} may have free type variables ϵ and ζ , the rest of $\hat{\chi}$ cannot.

Example In Figure 3, we show an example \mathbb{T} program demonstrating `call`, `jmp`, `ret`, and `halt`. The control flow, in Figure 4, shows the instructions causing jumps between basic blocks and the state of the relevant registers and stack at jump-time. In this diagram, l_2 and l_{2aux} are in the same component, while the rest are made up of distinct components that together make up the component f .

4. \mathbb{FT} Multi-Language

We present a minimal functional language \mathbb{F} , and then embed it and \mathbb{T} within a Matthews-Findler style multi-language. Particularly notable are the boundary translations for higher-order functions and code blocks. In §5, we design a logical relation with which we can show equivalence of programs that differ both structurally and algorithmically.

4.1 Functional Language: \mathbb{F}

In Figure 5 we present the syntax of \mathbb{F} , our simply-typed call-by-value functional language with isorecursive types, conditional branching, tuples, and base value integers and unit. The language is featureful enough to implement simple

Type τ	$::=$	$\alpha \mid \text{unit} \mid \text{int} \mid (\bar{\tau}) \rightarrow \tau \mid \mu\alpha.\tau \mid \langle \bar{\tau} \rangle$
Expression e	$::=$	$x \mid () \mid n \mid \text{tp} \mid \text{if0} \, t \, t \, t \mid \lambda(\bar{x}:\bar{\tau}).t \mid t \, \bar{t}$ $\text{fold}_{\mu\alpha.\tau} \, t \mid \text{unfold} \, t \mid \langle \bar{t} \rangle \mid \pi_i(\bar{t})$ where $p ::= + \mid - \mid *$
Value v	$::=$	$() \mid n \mid \lambda(\bar{x}:\bar{\tau}).t \mid \text{fold}_{\mu\alpha.\tau} \, v \mid \langle \bar{v} \rangle$
Evaluation ctxt E	$::=$	$[\cdot] \mid E \, p \, t \mid v \, p \, E \mid \text{if0} \, E \, t \, t \mid E \, \bar{t} \mid v \, \bar{v} \, E \, \bar{t}$ $\text{fold}_{\mu\alpha.\tau} \, E \mid \text{unfold} \, E \mid \langle \bar{v}, E, \bar{t} \rangle \mid \pi_i(E)$

Figure 5. \mathbb{F} Syntax

programs, while lacking certain expressiveness (like mutation) that we can add by way of the embedded assembly. The typing and operational semantics are standard and provided in our supplementary material.

4.2 Embedding \mathbb{T} in \mathbb{FT}

Syntax In Figure 6 we present the syntax of our multi-language \mathbb{FT} , which is largely made up of extensions to syntactic categories of either \mathbb{T} (Figure 1) or \mathbb{F} (Figure 5). Note that both expressions e and components e are components e in this language. Henceforth, when we refer to an \mathbb{F} or \mathbb{T} term we are referring to the terms that originated in that language, which can now of course include nested components of the other language. We add boundaries $\mathcal{TFT}e$ (\mathbb{T} inside, \mathbb{F} outside) and $\mathcal{FTT}e$ (\mathbb{F} inside, \mathbb{T} outside) to mediate between the languages. In both cases, the \mathbb{F} type τ directs the translation. In particular, the $\mathcal{TFT}e$ contains a \mathbb{T} component e with \mathbb{T} translated type $\tau^\mathcal{T}$, while the $\mathcal{FTT}e$ contains a \mathbb{F} expression e of type τ . Like Matthews-Findler [?], we reduce the component within the boundary to a value, after which we carry out a type-directed value translation using translations metafunctions $\mathcal{TFT}(\cdot)$ and $\mathcal{FTT}(\cdot)$:

$$\mathcal{TFT}e \mapsto^* \mathcal{TFT}v \mapsto \mathcal{FTT}(v)$$

To \mathbb{T} instructions ι , we add an `import` instruction to wrap the boundary and to specify what register the translated value should be placed in. The `import` instruction also specifies σ , the tail of the stack that should be protected while evaluating the \mathbb{F} expression e , which could in turn include \mathbb{T} code.

When translating \mathbb{T} code blocks into \mathbb{F} functions, we will need to instantiate the stack tail variable ζ on the \mathbb{T} code block. For this reason, we introduce the `protect` instruction, which specifies a stack prefix ϕ to leave visible and a type variable ζ to bind to the tail. We'll see the value translation later in the section.

While normal \mathbb{F} lambdas are embedded in the multi-language, they do not allow stack modification in embedded \mathbb{T} code. However, in some cases we explicitly will want to allow that sort of modification. For this reason, we introduce a new stack-modifying lambda term $\lambda_{\phi_o}^{\phi_i}(\bar{x}:\bar{\tau}).t$, which specifies the stack prefix ϕ_i it requires on the front of the stack when it is called, and the stack prefix ϕ_o that it will

Type τ	::=	$\dots \mid (\bar{\tau}) \xrightarrow{\phi_i \phi_o} \tau'$
Expression e	::=	$\dots \mid \mathcal{F}T e \mid \lambda_{\phi}^{\phi}(\bar{x}:\bar{\tau}).t \mid t \bar{t}'$
Return marker q	::=	$\dots \mid \text{out}$
Instruction ι	::=	$\dots \mid \text{import } r_d, {}^{\sigma}\mathcal{F}T e$
		$\text{protect } \phi, \zeta$
Stack prefix ϕ	::=	$\cdot \mid \tau :: \phi$
Stack typing σ	::=	$\phi :: \zeta \mid \phi :: \bullet$
Evaluation ctxt E	::=	$\dots \mid \mathcal{F}T E$
Evaluation ctxt E	::=	$\dots \mid (\text{import } r_d, {}^{\sigma}\mathcal{F}T E; I, \cdot)$
Type τ	::=	$\tau \mid \tau$
Component e	::=	$e \mid e$
Δ	::=	$\cdot \mid \Delta, \alpha \mid \Delta, \alpha \mid \Delta, \zeta \mid \Delta, \epsilon$
Evaluation ctxt E	::=	$E \mid E$

Figure 6. $\mathbb{F}\mathbb{T}$ Multi-Language Syntax

have replaced ϕ_i with upon return. Correspondingly, we introduce a new arrow type that captures that relationship. Note that the ordinary lambda can be seen as a special case when ϕ_i and ϕ_o are both the empty prefix \cdot , which corresponds to the entire stack being the protected tail.

We also add a new return marker, **out**, which is used for \mathbb{F} code, as \mathbb{F} follows normal expression-based evaluation and thus has no return continuation.

Type System The typing judgments for $\mathbb{F}\mathbb{T}$, for which we show a selection in Figure 7, include modified versions from both \mathbb{T} and \mathbb{F} judgments as well as rules for the new forms. Since this is a multi-language and not a compiler, the typing rules for \mathbb{T} must now include an \mathbb{F} environment Γ of free \mathbb{F} variables. Similarly, the typing rules for \mathbb{F} must now include all of the context needed by \mathbb{T} , since in order to type-check embedded assembly components we will need to know the current register (χ), stack (σ), and heap (Ψ) typings.

Most of these modifications are straightforward; we show the rule for \mathbb{F} application in Figure 7 as a representative. Note that the stack typings σ_i are threaded through the arguments according to evaluation order, as each one could include embedded \mathbb{T} code that modified the stack.

For the boundary term, $\mathcal{F}T e$, we require that the \mathbb{T} component e within the boundary be well typed under translation type τ^T and return marker $\text{end}\{\tau^T; \sigma'\}$, which corresponds to the inner assembly halting with a value of type τ^T . In that case, the boundary term is well typed under τ at the **out** return marker that corresponds to \mathbb{F} code. Note that the boundary makes no restriction on modification of the stack. Also in the figure is the typing rule for the stack-modifying lambda term, which is an ordinary lambda typing rule with the exception that it types under stacks with the given prefixes ϕ_i and ϕ_o and abstract tails ζ .

$\Psi; \Delta; \Gamma; \chi; \sigma; q \vdash e: \tau; \sigma'$
$\Psi; \Delta; \Gamma; \chi; \sigma; \text{out} \vdash t: (\tau_1 \dots \tau_n) \rightarrow \tau'; \sigma_0$
$\Psi; \Delta; \Gamma; \chi; \sigma_{i-1}; \text{out} \vdash t_i: \tau_i; \sigma_i$
$\Psi; \Delta; \Gamma; \chi; \sigma; \text{out} \vdash t \tau_1 \dots \tau_n: \tau'; \sigma_n$
$\Psi; \Delta; \Gamma; \cdot; \sigma; \text{end}\{\tau^T; \sigma'\} \vdash e: \tau^T; \sigma'$
$\Psi; \Delta; \Gamma; \chi; \sigma; \text{out} \vdash \mathcal{F}T e: \tau; \sigma'$
$\Psi; \Delta, \zeta; \Gamma; \bar{x}:\bar{\tau}; \chi; \phi_i :: \zeta; \text{out} \vdash t: \tau'; \phi_o :: \zeta$
$\Psi; \Delta; \Gamma; \chi; \sigma; \text{out} \vdash \lambda_{\phi_o}^{\phi_i}(\bar{x}:\bar{\tau}).t: (\bar{\tau}) \xrightarrow{\phi_i \phi_o} \tau'; \sigma$
$\Psi; \Delta; \Gamma; \chi; \sigma; q \vdash \iota \Rightarrow \Delta'; \chi'; \sigma'; q'$
$\sigma = \phi :: \sigma_0 \quad \sigma' = \phi :: \zeta$
$\Psi; \Delta; \Gamma; \chi; \sigma; q \vdash \text{protect } \phi, \zeta \Rightarrow \Delta, \zeta; \chi; \sigma'; q$
$\sigma = \tau_0 :: \dots :: \tau_j :: \sigma_0 \quad \sigma' = \tau'_0 :: \dots :: \tau'_k :: \sigma_0$
$\Psi; \Delta, \zeta; \Gamma; \chi; (\tau_0 :: \dots :: \tau_j :: \zeta); \text{out} \vdash e: \tau; (\tau'_0 :: \dots :: \tau'_k :: \zeta)$
$q = i > j \text{ or } q = \text{end}\{\hat{\tau}; \hat{\sigma}\}$
$\Psi; \Delta; \Gamma; \chi; \sigma; q \vdash \text{import } r_d, {}^{\sigma_0}\mathcal{F}T e \Rightarrow \Delta; (r_d: \tau^T); \sigma'; \text{inc}(q, k-j)$

Figure 7. Selected $\mathbb{F}\mathbb{T}$ Typing Rules

As described above, we add two new \mathbb{T} instructions. The **protect** instruction is used to abstract the tail of the stack, which we can see in the transformation of the stack $\phi :: \sigma_0$ before **protect** into $\phi :: \zeta$ after, where ζ is a new type variable introduced to the type environment. Note that there is no way to undo this; it will last until the end of the current \mathbb{T} component.

The other new instruction is the \mathbb{T} boundary instruction **import**. Ignoring stacks, the rule is quite simple: it takes an \mathbb{F} term e of type τ , well typed under the **out** return marker, and translates it to type τ^T , storing the result in register r_d . This story is complicated by the handling of stacks, as it is important for **import** instructions to be able to restrict what portion of the stack the inner code can modify. In particular, since the \mathbb{F} code does not have the same return marker q , we must be sure that q cannot be clobbered by \mathbb{T} code embedded in e . To do this, we specify the portion of the stack σ_0 that is abstracted as ζ in e , and ensure that either q is stored in that stack tail or it is the halting marker. Finally, since the front of the stack could grow or shrink to k entries, if q were a stack index i we increment it by $k - j$ using the metafunction inc , which otherwise is identity.

Operational Semantics The operational semantics for boundary terms, shown in Figure 8, translate values using the type-directed metafunctions $\mathcal{F}\mathbb{T}(\cdot)$ (\mathbb{T} inside, \mathbb{F} outside) and $\mathcal{T}\mathbb{F}(\cdot)$ (\mathbb{F} inside, \mathbb{T} outside).

$$\begin{aligned}
&\langle \mathbf{M} \mid E[\mathcal{TFT}(\text{ret end}\{\tau^{\mathcal{T}}; \sigma\} \{r\}, \cdot)] \rangle \\
&\quad \mapsto \langle \mathbf{M}' \mid E[\mathbf{v}] \rangle \quad \text{if } \mathcal{TFT}(\mathbf{M.R}(r), \mathbf{M}) = (\mathbf{v}, \mathbf{M}') \\
&\langle \mathbf{M} \mid E[\text{import } r_d, \sigma' \mathcal{TFT}^{\mathcal{T}} \mathbf{v}; \mathbf{I}] \rangle \\
&\quad \mapsto \langle \mathbf{M}' \mid E[\text{mv } r_d, \mathbf{w}; \mathbf{I}] \rangle \text{ if } \mathcal{TFT}^{\mathcal{T}}(\mathbf{v}, \mathbf{M}) = (\mathbf{w}, \mathbf{M}')
\end{aligned}$$

Figure 8. \mathcal{TFT} Operational Semantics: Language Boundaries

$$\begin{aligned}
&\alpha^{\mathcal{T}} = \alpha \\
&\text{unit}^{\mathcal{T}} = \text{unit} \quad \mu\alpha.\tau^{\mathcal{T}} = \mu\alpha.(\tau^{\mathcal{T}}) \\
&\text{int}^{\mathcal{T}} = \text{int} \quad \langle \tau_1, \dots, \tau_n \rangle^{\mathcal{T}} = \text{box} \langle \tau_1^{\mathcal{T}}, \dots, \tau_n^{\mathcal{T}} \rangle \\
&(\tau_1, \dots, \tau_n) \rightarrow \tau'^{\mathcal{T}} = \\
&\quad \text{box } \forall[\zeta, \epsilon]. \{ra: \text{box } \forall[].\{r1: \tau'^{\mathcal{T}}; \zeta\}^{\epsilon}; \sigma'\}^{ra} \\
&\quad \text{where } \sigma' = \tau_n^{\mathcal{T}} :: \dots :: \tau_1^{\mathcal{T}} :: \zeta \\
&(\tau_1, \dots, \tau_n) \xrightarrow{\phi_i: \phi_o} \tau'^{\mathcal{T}} = \\
&\quad \text{box } \forall[\zeta, \epsilon]. \{ra: \text{box } \forall[].\{r1: \tau'^{\mathcal{T}}; \phi_o :: \zeta\}^{\epsilon}; \sigma'\}^{ra} \\
&\quad \text{where } \sigma' = \tau_n^{\mathcal{T}} :: \dots :: \tau_1^{\mathcal{T}} :: \phi_i :: \zeta
\end{aligned}$$

Figure 9. \mathcal{TFT} Boundary Type Translation

Figure 9 contains the type translation guiding these meta-functions. Note that \mathbb{F} tuples are translated to immutable references to \mathbb{T} heap tuples. The most complex transformation is for function types, which are translated into code blocks that pass arguments on the stack and follow the calling convention described in §3 where return continuations can be instantiated alternately by \mathbb{T} or \mathbb{F} callers.

We show the value translations in Figure 10, eliding only the stack-modifying lambda, which is similar to the lambda shown. The most significant translations are between \mathbb{T} code blocks and \mathbb{F} functions. In particular, we must translate between variable representations and calling conventions — this means the arguments are passed on the stack, and a return continuation must be in register ra . Finally, we must translate the arguments themselves, and translate the return value back, cleaning up temporary stack values.

Critically, when translating an \mathbb{F} function to a \mathbb{T} code block, we must protect the return continuation, since embedded assembly blocks within the body of the function could write to register ra . To do that, we store it on the stack and protect the tail. In the stack-modifying lambda case, this is complicated slightly by needing to re-arrange the stack to put the protected value past the exposed stack prefix ϕ_i .

Then, to evaluate the \mathbb{F} function, we load each argument off of the stack, translate it to \mathbb{F} , apply the function, and import the returned value back to \mathbb{T} . After doing this, we load the return continuation off of the stack, clear the arguments according to the calling convention and return. Note that in the stack-modifying lambda case, we have to be careful to clear the arguments but keep the output prefix ϕ_o .

$$\begin{aligned}
&\mathcal{TFT}^{\text{int}}(\mathbf{n}, \mathbf{M}) = (\mathbf{n}, \mathbf{M}) \\
&\mathcal{TFT}^{\mu\alpha.\tau}(\text{fold}_{\mu\alpha.\tau} \mathbf{v}, \mathbf{M}) = (\text{fold}_{\mu\alpha.\tau} \mathbf{v}, \mathbf{M}') \\
&\quad \text{where } \mathcal{TFT}^{\tau[\mu\alpha.\tau/\alpha]}(\mathbf{v}, \mathbf{M}) = (\mathbf{v}, \mathbf{M}') \\
&\mathcal{TFT}^{\langle \tau_1, \dots, \tau_n \rangle}(\langle \mathbf{v}_0, \dots, \mathbf{v}_n \rangle, \mathbf{M}) = \\
&\quad (\ell, (\mathbf{M}_{n+1}, \ell \mapsto \langle \mathbf{w}_0, \dots, \mathbf{w}_n \rangle)) \\
&\quad \text{where } \mathbf{M}_0 = \mathbf{M}, \text{ and } \mathcal{TFT}^{\tau}(\mathbf{v}_i, \mathbf{M}_i) = (\mathbf{w}_i, \mathbf{M}_{i+1}) \\
&\mathcal{TFT}^{\text{unit}}(\cdot, \mathbf{M}) = (\cdot, \mathbf{M}) \\
&\mathcal{TFT}^{\langle \tau \rangle \rightarrow \tau'}(\lambda(\bar{x}:\bar{\tau}).\mathbf{t}, \mathbf{M}) = (\ell, (\mathbf{M}, \ell \mapsto \mathbf{h})) \\
&\quad \text{where } \mathbf{h} = \text{code}[\zeta, \epsilon]\{ra: \forall[].\{r1: \tau'^{\mathcal{T}}; \zeta\}^{\epsilon}; \tau^{\mathcal{T}} :: \zeta\}^{ra}. \\
&\quad \text{salloc } 1; \text{sst } 0, ra; \text{import } r_1, \zeta \mathcal{TFT}^{\tau'} \mathbf{e}; \\
&\quad \text{sld } ra, 0; \text{sfree } n+1; \text{ret } ra \{r1\} \\
&\quad \mathbf{e} = (\lambda(\bar{x}:\bar{\tau}).\mathbf{t})^{\tau \mathcal{TFT}}(\text{sld } r_1, n+1-i; \\
&\quad \quad \text{ret end}\{\tau^{\mathcal{T}}; \sigma\} \{r1\}, \cdot) \\
&\quad \sigma = \forall[].\{r1: \tau'^{\mathcal{T}}; \zeta\}^{\epsilon} :: \tau^{\mathcal{T}} :: \zeta \\
&\text{unit}\mathcal{TFT}(\cdot, \mathbf{M}) = (\cdot, \mathbf{M}) \\
&\text{int}\mathcal{TFT}(\mathbf{n}, \mathbf{M}) = (\mathbf{n}, \mathbf{M}) \\
&\mu\alpha.\tau\mathcal{TFT}(\text{fold}_{\mu\alpha.\tau} \mathbf{w}, \mathbf{M}) = (\text{fold}_{\mu\alpha.\tau} \mathbf{v}, \mathbf{M}') \\
&\quad \text{where } \tau[\mu\alpha.\tau/\alpha]\mathcal{TFT}(\mathbf{w}, \mathbf{M}) = (\mathbf{v}, \mathbf{M}') \\
&\langle \tau_0, \dots, \tau_n \rangle \mathcal{TFT}(\ell, \mathbf{M}) = (\langle \mathbf{v}_0, \dots, \mathbf{v}_n \rangle, \mathbf{M}_{n+1}) \\
&\quad \text{where } \mathbf{M}(\ell) = \langle \mathbf{w}_0, \dots, \mathbf{w}_n \rangle, \\
&\quad \mathbf{M}_0 = \mathbf{M}, \text{ and } \tau\mathcal{TFT}(\mathbf{w}_i, \mathbf{M}_i) = (\mathbf{v}_i, \mathbf{M}_{i+1}) \\
&\langle \tau_n \rangle \rightarrow \tau' \mathcal{TFT}(\mathbf{w}, \mathbf{M}) = (\mathbf{v}, (\mathbf{M}, \ell_{\text{end}} \mapsto \mathbf{h}_{\text{end}})) \\
&\quad \text{where } \mathbf{v} = \lambda(\bar{x}_n:\bar{\tau}_n).\tau' \mathcal{TFT}(\text{protect } \cdot, \zeta; \text{import } r_1, \zeta \mathcal{TFT}^{\tau_1} \mathbf{x}_1; \\
&\quad \quad \text{salloc } 1; \text{sst } 0, r_1; \dots; \\
&\quad \quad \text{import } r_1, \zeta \mathcal{TFT}^{\tau_n} \mathbf{x}_n; \text{salloc } 1; \text{sst } 0, r_1; \\
&\quad \quad \text{mv } ra, \ell_{\text{end}}[\zeta]; \text{jmp } w[\zeta][\text{end}\{\tau'^{\mathcal{T}}; \zeta\}, \cdot) \\
&\quad \mathbf{h}_{\text{end}} = \text{code}[\zeta]\{r1: \tau'^{\mathcal{T}}; \zeta\}^{\epsilon} \text{end}\{\tau'^{\mathcal{T}}; \zeta\}. \\
&\quad \quad \text{ret end}\{\tau'^{\mathcal{T}}; \zeta\} \{r1\}
\end{aligned}$$

Figure 10. \mathcal{TFT} Boundary Value Translation

Example In Figure 11 we show a program where an \mathbb{F} component \mathbf{e} calls a \mathbb{T} component ℓ passing it an \mathbb{F} function \mathbf{g} that in turn accepts a function (\mathbf{h}) and calls it on $\mathbf{1}$. The \mathbb{T} component ℓ can then pass a \mathbb{T} code pointer ℓ_{aux} as the argument to \mathbf{g} , returning the value produced back through \mathbb{T} and out to \mathbb{F} .

We show a control-flow diagram for this example in Figure 12, where arrows in \mathbb{F} boxes correspond to argument passing and return values, whereas within \mathbb{T} boxes arrows correspond to jumps or halt, as in Figure 4. In the diagram, we color the blocks gray that were introduced by the multi-language value translation: ℓ_{hret} and ℓ_{ret} . Every \mathbb{T} box will terminate with a **halt**, at which point control passes back to \mathbb{F} . When crossing the language boundary an \mathbb{F} application corresponds to a \mathbb{T} **call** and vice versa, whereas


```

 $\tau = ((\text{int}) \rightarrow \text{int}) \rightarrow \text{int}$ 
 $g = \lambda(h: (\text{int}) \rightarrow \text{int}).h\ 1$ 
 $e = (\text{int}^{\mathcal{FT}}(\text{mv } r1, \ell; \text{halt } (\tau) \rightarrow \text{int}^{\mathcal{T}}, \bullet \{r1\}, H))\ g$ 
 $H(\ell) = \text{code}[\zeta, \epsilon]\{\text{ra}: \forall[].\{r1: \text{int}^{\mathcal{T}}; \zeta\}^{\epsilon}; \tau^{\mathcal{T}} :: \zeta\}^{\text{ra}}.$ 
 $\quad \text{sld } r1, 0; \text{salloc } 1; \text{mv } r2, \ell_h; \text{sst } 0, r2; \text{sst } 1, \text{ra};$ 
 $\quad \text{mv } \text{ra}, \ell_{\text{gret}}[\zeta, \epsilon]; \text{call } r1 \{\forall[].\{r1: \text{int}^{\mathcal{T}}; \zeta\}^{\epsilon} :: \zeta, 0\}$ 
 $H(\ell_h) = \text{code}[\zeta, \epsilon]\{\text{ra}: \forall[].\{r1: \text{int}^{\mathcal{T}}; \zeta\}^{\epsilon}; \text{int}^{\mathcal{T}} :: \zeta\}^{\text{ra}}.$ 
 $\quad \text{sld } r1, 0; \text{sfree } 1; \text{mul } r1, r1, 2; \text{ret } \text{ra} \{r1\}$ 
 $H(\ell_{\text{gret}}) = \text{code}[\zeta, \epsilon]\{r1: \text{int}; \forall[].\{r1: \text{int}^{\mathcal{T}}; \zeta\}^{\epsilon} :: \zeta\}^0.$ 
 $\quad \text{sld } \text{ra}, 0; \text{sfree } 1; \text{ret } \text{ra} \{r1\}$ 

```

Figure 11. \mathcal{FT} Example: Higher-order

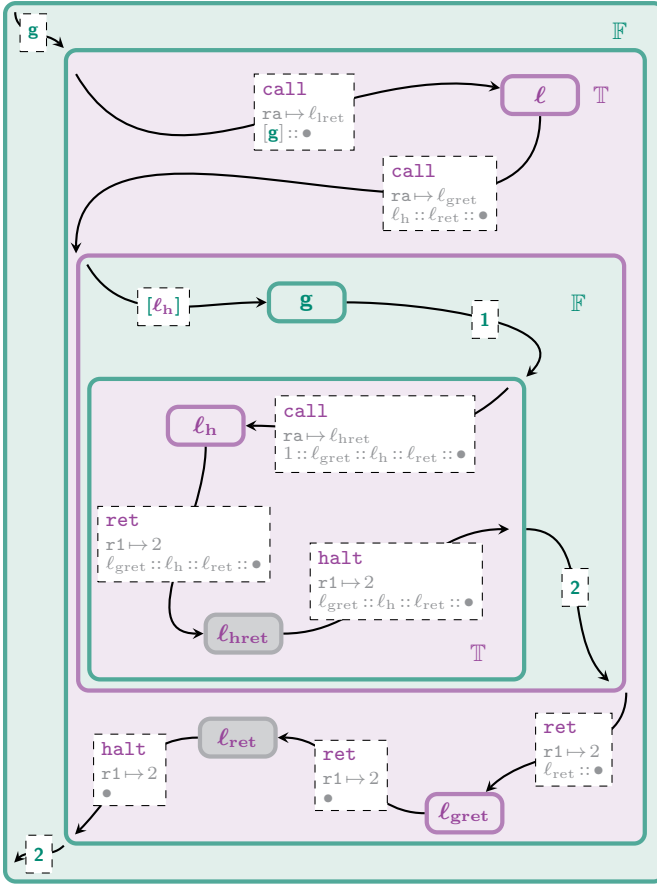


Figure 12. \mathcal{FT} Control Flow: Higher-order (Fig. 11)

while an \mathcal{F} return corresponds to a \mathcal{T} `ret`, a \mathcal{T} `halt` corresponds to an \mathcal{F} return.

This short example demonstrates the ability of our model to reflect arbitrary nesting across language boundaries, as can happen when modeling JIT compilation. In this case, we imagine that the portion of code containing ℓ and ℓ_h has been compiled by the JIT, but e and g have not.

5. Logical Relation for \mathcal{FT}

In order to reason about program equivalence in \mathcal{FT} , we design a step-indexed Kripke logical relation for our language. We take our essential design from Dreyer *et al.* [?] and Ahmed *et al.* [?], where the Kripke worlds contain *islands* with state-transition systems to accommodate mutations to the heap, registers, and stack. From their model, we get the ability to reason about equivalences dependent on hidden mutable state. We do not go into detail on this in this paper, but this power is the reason that we base our work on their well-established model. In this section, we explore the ways that we had to adapt the model to the setting of \mathcal{FT} .

In our logical relation, for which we show the closed relations in Figure 13, we have three value relations: $\mathcal{V}[\tau]\rho$, $\mathcal{W}[\tau]\rho$, and $\mathcal{HV}[\psi]\rho$, which correspond to the three types of values that exist in \mathcal{FT} : high-level values, low-level word-sized values, and low-level heap values. In these relations, as usual, ρ is a relational substitution for type variables. Further, with the exception of contexts in the \mathcal{K} relation, all of our relations are built out of well-typed terms, though we elide that requirement in these figures.

As a Kripke logical relation, relatedness of values depends on the state of the a world W . Some values are related irrespective of world state; for example, $(W, \mathbf{n}, \mathbf{n}) \in \mathcal{W}[\text{int}]\rho$ for any W . However, the structure of the world captures key semantic properties about the stack, heap, and registers in a sequence of *islands* that describe the current state of memories. Each island expresses invariants on certain parts of memory by encoding a state-transition system and a memory relation that establishes which pairs of memories are related in each state.

Since our logical relation is step-indexed, our worlds have an index k , which conveys that the semantic equivalence of the relation is true for at least k steps, but no information is known beyond that. This allows us to avoid circularity when dealing with recursive types, as we can induct on the step index rather than the structure of the expanding type.

$W' \sqsubseteq W$ says W' is a future world of W ; to reach it, we may have consumed steps (lowering k), allocated additional memory in new islands, or made transitions in islands.

While most of our logical relation follows prior work [?], here we focus on the parts that are novel. In particular, since we are dealing with assembly we must reason about the semantic equivalence of code blocks at code-pointer type.

Our code pointer logical relation (Figure 15) is like a function logical relation, in that given related inputs, it should produce related outputs. Inputs, in this case, are registers and the stack, for which we have the requirement that in a future world W' with closing type substitution ρ^* , $\text{curr-R}(W') \subseteq \mathcal{R}[\chi]\rho'$ and $\text{curr-S}(W') \subseteq \mathcal{S}[\sigma]\rho'$. This means that the current register files and stacks in world W' are related at register file typing χ and stack typing σ respectively. Related register files map registers to related values,

Statement	Meaning
$(W, \mathbf{v}_1, \mathbf{v}_2) \in \mathcal{V}[\![\tau]\!]\rho$	\mathbf{v}_1 and \mathbf{v}_2 are related \mathbb{F} values at type τ in world W under type substitution ρ
$(W, \mathbf{w}_1, \mathbf{w}_2) \in \mathcal{W}[\![\tau]\!]\rho$	\mathbf{w}_1 and \mathbf{w}_2 are related \mathbb{T} word values at type τ in world W under type substitution ρ
$(W, \mathbf{h}_1, \mathbf{h}_2) \in \mathcal{H}\mathcal{V}[\![\psi]\!]\rho$	\mathbf{h}_1 and \mathbf{h}_2 are related \mathbb{T} heap values at type ψ in world W under type substitution ρ
$(W, e_1, e_2) \in \mathcal{O}$	e_1 and e_2 run with memories related at W , either both terminate or are both running after $W.k$ steps
$(W, E_1, E_2) \in \mathcal{K}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho$	E_1 and E_2 are related continuations, so given appropriately related values at type τ , they are in \mathcal{O}
$(W, e_1, e_2) \in \mathcal{E}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho$	e_1 and e_2 are related expressions, so given appropriate related continuations, they are in \mathcal{O}

Figure 13. $\mathbb{F}\mathbb{T}$ Logical Relation: Closed Values and Terms

$$\begin{aligned}
\mathcal{E}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho &= \{ (W, e_1, e_2) \mid \forall E_1, E_2. (W, E_1, E_2) \in \mathcal{K}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho \implies (W, E_1[e_1], E_2[e_2]) \in \mathcal{O} \} \\
\mathcal{K}[\![\text{out} \vdash \tau; \sigma]\!]\rho &= \{ (W, E_1, E_2) \mid \forall W', \mathbf{v}_1, \mathbf{v}_2. W' \sqsupseteq_{\text{pub}} W \wedge (W', \mathbf{v}_1, \mathbf{v}_2) \in \mathcal{V}[\![\tau]\!]\rho \wedge \text{curr-S}(W') \in \mathcal{S}[\![\sigma]\!]\rho \\
&\implies (W', E_1[\mathbf{v}_1], E_2[\mathbf{v}_2]) \in \mathcal{O} \} \\
\mathcal{K}[\![\text{end}\{\tau; \sigma\} \vdash \tau; \sigma]\!]\rho &= \{ (W, E_1, E_2) \mid \forall W', \mathbf{r}_1, \mathbf{r}_2. W' \sqsupseteq_{\text{pub}} W \wedge \\
&(\triangleright W', W'.\mathbf{R}_1(\mathbf{r}_1), W'.\mathbf{R}_2(\mathbf{r}_2)) \in \mathcal{W}[\![\tau]\!]\rho \wedge \text{curr-S}(W') \in \mathcal{S}[\![\sigma]\!]\rho \\
&\implies (W', E_1[(\text{halt } \rho_1(\tau), \rho_1(\sigma) \{ \mathbf{r}_1 \}, \cdot)], E_2[(\text{halt } \rho_2(\tau), \rho_2(\sigma) \{ \mathbf{r}_2 \}, \cdot)]) \in \mathcal{O} \} \\
\mathcal{K}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho &= \{ (W, E_1, E_2) \mid (\mathbf{q} = \mathbf{r} \vee \mathbf{q} = \mathbf{i}) \wedge \forall W', \mathbf{q}', \mathbf{r}_1, \mathbf{r}_2. W' \sqsupseteq_{\text{pub}} W \wedge \\
&(\exists \mathbf{r}. \mathbf{q}' = \mathbf{r} \wedge \text{ret-addr}_1(W, \rho_1(\mathbf{q})) = W'.\mathbf{R}_1(\mathbf{r}) \wedge \text{ret-addr}_2(W, \rho_2(\mathbf{q})) = W'.\mathbf{R}_2(\mathbf{r}) \wedge \\
&\text{ret-reg}_1(W', \mathbf{r}) = \mathbf{r}_1 \wedge \text{ret-reg}_2(W', \mathbf{r}) = \mathbf{r}_2) \wedge \\
&(\triangleright W', W'.\mathbf{R}_1(\mathbf{r}_1), W'.\mathbf{R}_2(\mathbf{r}_2)) \in \mathcal{W}[\![\tau]\!]\rho \wedge \text{curr-S}(W') \in \mathcal{S}[\![\sigma]\!]\rho \\
&\implies (W', E_1[(\text{ret } \rho_1(\mathbf{q}') \{ \mathbf{r}_1 \}, \cdot)], E_2[(\text{ret } \rho_2(\mathbf{q}') \{ \mathbf{r}_2 \}, \cdot)]) \in \mathcal{O} \} \\
\text{ret-addr}_j(W, \mathbf{r}) &= W.\mathbf{R}_j(\mathbf{r}) \quad \text{ret-addr}_j(W, \mathbf{i}) = W.\mathbf{S}_j(\mathbf{i}) \quad \text{ret-reg}_j(W, \mathbf{r}) = \mathbf{r}' \text{ if } W.\chi_j(\mathbf{r}) = \text{box } \forall \mathbf{r}'. \{ \mathbf{r}' : \tau; \sigma' \}^a \\
\Psi; \Delta; \Gamma; \chi; \sigma; \mathbf{q} \vdash e_1 \approx e_2 : \tau; \sigma' &\stackrel{\text{def}}{=} \forall W, \gamma, \rho. W \in \mathcal{H}[\![\Psi]\!] \wedge \rho \in \mathcal{D}[\![\Delta]\!] \wedge (W, \gamma) \in \mathcal{G}[\![\Gamma]\!]\rho \wedge \text{curr-R}(W) \in \mathcal{R}[\![\chi]\!]\rho \wedge \\
&\text{curr-S}(W) \in \mathcal{S}[\![\sigma]\!]\rho \implies (W, \rho_1(\gamma_1(e_1)), \rho_2(\gamma_2(e_2))) \in \mathcal{E}[\![\mathbf{q} \vdash \tau; \sigma']\!]\rho
\end{aligned}$$

Figure 14. $\mathbb{F}\mathbb{T}$: Component and Continuation Relations and Equivalence of Open Terms

$$\begin{aligned}
\mathcal{H}\mathcal{V}[\![\forall[\Delta].\{\chi; \sigma\}^a]\!]\rho &= \\
&\{ (W, \text{code}[\Delta]\{\rho_1(\chi); \rho_1(\sigma)\}^{\rho_1(\mathbf{q})}. \mathbf{I}_1, \\
&\quad \text{code}[\Delta]\{\rho_2(\chi); \rho_2(\sigma)\}^{\rho_2(\mathbf{q})}. \mathbf{I}_2) \mid \\
&\quad \forall W' \sqsupseteq W. \forall \rho^* \in \mathcal{D}[\![\Delta]\!]. \forall \tau, \sigma'. \\
&\quad \text{let } \rho' = \rho \cup \rho^* \text{ in } \tau; \sigma' =_{\rho'} \text{ret-type}(\mathbf{q}, \chi, \sigma) \wedge \\
&\quad \text{curr-R}(W') \in \mathcal{R}[\![\chi]\!]\rho' \wedge \text{curr-S}(W') \in \mathcal{S}[\![\sigma]\!]\rho' \\
&\implies (W', (\rho_1^*(\mathbf{I}_1), \cdot), (\rho_2^*(\mathbf{I}_2), \cdot)) \in \mathcal{E}[\![\mathbf{q} \vdash \tau; \sigma']\!]\rho' \} \\
\tau; \sigma' =_{\rho} \text{ret-type}(\mathbf{q}, \chi, \sigma) &\stackrel{\text{def}}{=} \\
\rho_i(\tau); \rho_i(\sigma') &= \text{ret-type}(\rho_i(\mathbf{q}), \rho_i(\chi), \rho_i(\sigma)), \text{ for } i \in 1, 2
\end{aligned}$$

Figure 15. $\mathbb{F}\mathbb{T}$ Logical Relation: Code Block

and related stacks are made up of related values. Stacks are related at ζ if they are related by relational substitution ρ' .

Once we have related inputs, the logical relation should specify that applying the arguments produces related outputs expressions. Since the arguments are present in the registers and on the stack, we simply state that the instruction sequences \mathbf{I}_1 and \mathbf{I}_2 , with empty heap fragments, are related

components in the \mathcal{E} relation under those conditions, relying critically on the return marker \mathbf{q} to determine the return type τ and resulting stack σ' .

The logical relation \mathcal{E} for components has three formal parameters: \mathbf{q} , τ , and σ' . The return marker \mathbf{q} , says where the expression is returning to, as we described in §3. The return type τ is the type of value that is passed to the return continuation in \mathbf{q} , which is critical in order to reason about equivalences, because if expressions don't even produce the same type of value they can't possibly be equivalent. This type comes from the ret-type metafunction whose definition is in Figure 2. The output stack type σ' is also, in a sense, part of the return value, and it similarly is derived from the return marker by the metafunctions.

The component relation $\mathcal{E}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho$ and relation for evaluation contexts $\mathcal{K}[\![\mathbf{q} \vdash \tau; \sigma]\!]\rho$ are tightly connected, as is standard for logical relations based on biorthogonality. In typical biorthogonal presentations, the definitions would be:

$$\begin{aligned}
\mathcal{K}[\![\tau]\!] &= \{ (W, E_1, E_2) \mid \forall W'. W' \sqsupseteq W \wedge (W', v_1, v_2) \in \mathcal{V}[\![\tau]\!] \\
&\implies (W', E_1[v_1], E_2[v_2]) \in \mathcal{O} \} \\
\mathcal{E}[\![\tau]\!] &= \{ (W, e_1, e_2) \mid \forall E_1 E_2. (W, E_1, E_2) \in \mathcal{K}[\![\tau]\!] \\
&\implies (W, E_1[e_1], E_2[e_2]) \in \mathcal{O} \}
\end{aligned}$$

Which states that continuations E_1 and E_2 accepting type τ related at world W must be such that, given any future world W' and τ values, plugging in the values results in related observations. In turn, expressions e_1 and e_2 of type τ related at world W must be such that, given related continuations E_1 and E_2 , $E_1[e_1]$ and $E_2[e_2]$ are observationally equivalent. Note how the reduction of e_1 and e_2 to values is central, since the definition of E_1 and E_2 tells you only that given related values they produce related observation. This reduction is normally captured in “monadic bind” lemmas.

Our definitions, in Figure 14, are more involved, but follow exactly this pattern. Our relation \mathcal{E} only differs from the standard one in that the type of a component involves a return marker \mathbf{q} and output stack type σ .

The continuation relation \mathcal{K} has three cases for different the return marker \mathbf{q} . The case for **out**, which corresponds to our functional terms, is nearly identical to the idealized case shown above, only differing by ensuring $\text{curr-S}(W') \subseteq S[\![\sigma]\!] \rho$, which means that at the point we are plugging in the values \mathbf{v}_1 and \mathbf{v}_2 the stacks are related at the right type σ .

The \mathcal{K} relation for **end** $\{\tau; \sigma\}$ is similar, but since this is \mathbb{T} code, return values are stored in registers \mathbf{r}_i and the “value” being plugged in is the **halt** instruction.

The third case, when the return marker is a register \mathbf{r} or a stack position \mathbf{i} , is more involved, though the overall meaning is still the same as the other cases: in the future, we will have a value to pass and will plug it into the hole to get related observations. First, we note that while at points the return marker can be a stack index \mathbf{i} , when we actually return to the continuation the return marker must be stored in a register \mathbf{q}' . We require, however, that the code block being pointed to by \mathbf{q} is the same as what is pointed to by \mathbf{q}' . Next, we find the register \mathbf{r}_i where the return value will be passed, and ensure that these contain related values. Finally, we check that the stacks are related at the right type with $\text{curr-S}(W') \subseteq S[\![\sigma]\!] \rho$, before saying that plugging in the return must yield related observations.

Having described how closed terms are related, we lift this to open terms with \approx , shown at the bottom of Figure 14. We choose appropriate closing type and term substitutions, where $\mathcal{G}[\![\Gamma]\!] \rho$ is a relational substitution mapping \mathbb{F} variables to related \mathbb{F} values, and then state the equivalence after closing with these substitutions.

We have proven that the logical relation is sound and complete with respect to \mathbb{FT} contextual equivalence (see supplementary material).

Theorem 5.1 (Fundamental Property)

If $\Psi; \Delta; \Gamma; \chi; \sigma; \mathbf{q} \vdash e : \tau; \sigma'$ then
 $\Psi; \Delta; \Gamma; \chi; \sigma; \mathbf{q} \vdash e \approx e : \tau; \sigma'.$

As usual, we prove compatibility lemmas corresponding to typing rules, after which the fundamental property follows as a corollary. While the none of the compatibility lemmas for \mathbb{T} instructions are trivial, the most interesting one is for

```

f1      = λ(x : int).(int) → intFT(protect ·, ζ; mv r1, ℓ;
                                     halt (int) → intT, ζ {r1}, H1) x
H1(ℓ)    = code[ζ, ε]{ra: ∇[] . {r1: intT; ζ}^ε; intT :: ζ}^ra.
          sld r1, 0; add r1, r1, 1; add r1, r1, 1;
          sfree 1; ret ra {r1}
f2      = λ(x : int).(int) → intFT(protect ·, ζ; mv r1, ℓ;
                                     halt intT, ζ {r1}, H2) x
H2(ℓ)    = code[ζ, ε]{ra: ∇[] . {r1: intT; ζ}^ε; intT :: ζ}^ra.
          sld r1, 0; add r1, r1, 1; sst 0, r1; jmp ℓ' [ζ][ε]
H2(ℓ')   = code[ζ, ε]{ra: ∇[] . {r1: intT; ζ}^ε; intT :: ζ}^ra.
          sld r1, 0; add r1, r1, 1; sfree 1; ret ra {r1}

```

Figure 16. \mathbb{FT} Example: Different Number of Basic Blocks

```

factF    = λ(x : int).(F (foldμα.(α) → int F)) x
F         = λ(f : μα.(α) → int).λ(x : int).
          if0 x 1 ((unfold f) f) (x - 1) * x
factT    = λ(x : int).(int) → intFT(ℓfact, H) x
H(ℓfact)  = code[ζ, ε]{ra: ∇[] . {r1: intT; ζ}^ε; intT :: ζ}^ra.
          sld rn, 0; mv rr, 1; bnz rn, ℓloop [ζ][ε];
          sfree 1; ret ra {rr}
H(ℓloop)  = code[ζ, ε]{rn: int, rr: int,
                    ra: ∇[] . {r1: intT; ζ}^ε; intT :: ζ}^ra.
          mul rr, rr, rn; sub rn, rn, 1;
          bnz rn, ℓloop [ζ][ε]; sfree 1; ret ra {rr}

```

Figure 17. \mathbb{FT} Example: Factorial Two Different Ways

the **call** instruction. In particular, **call** must ensure that the code that it is jumping to eventually returns, even while the target component could make nested calls. This relies on the the target component return marker ensuring that control will eventually pass to the original return continuation.

Theorem 5.2 (LR Sound & Complete wrt Ctx Equiv)

$\Psi; \Delta; \Gamma; \chi; \sigma; \mathbf{q} \vdash e_1 \approx e_2 : \tau; \sigma'$ if and only if
 $\Psi; \Delta; \Gamma; \chi; \sigma; \mathbf{q} \vdash e_1 \approx^{ctx} e_2 : \tau; \sigma'.$

5.1 Example Equivalences

In Figure 16, we show two programs that differ in the number of basic blocks that they use to carry out the same computation: adding two to a number and returning it. This example demonstrates our ability to reason over differences in internal jumps, which critically depends on the return markers explained in §3. We are able to show these two examples equivalent at type $(\text{int}) \rightarrow \text{int}$ using the logical relation. The elided proofs are included in the supplementary material.

In Figure 17, we show another small example. We present two implementations of the factorial function. The **fact_F** is a standard recursive functional implementation in the setting of isorecursive types. We apply the function template **F** to a folded version of itself and the argument **x**, and in the body we check if the **x** is **0**, in which case we return 1, and otherwise we unfold the first argument, call in with **x - 1**, and multiply the result by **x**. This clearly produces the result for **x ≥ 0**, and also clearly diverges for negative arguments.

fact_T is an imperative factorial, which uses mutates registers to compute the result. It has two basic blocks, ℓ_{fact} and ℓ_{loop} . The first, which is translated to **F** and called with argument **x**, loads the argument **n** (translated from **n**) into register **rn**, stores **1** in the result register **rr**, and then checks if **rn** is **0**. If so, we clear the argument off the stack and return. Otherwise, we jump to ℓ_{loop} . This multiplies the result by **rn**, subtracts one from **rn**, and makes the same check if **rn** is zero. If so, we do the same cleanup and return, and otherwise we jump to the beginning of ℓ_{loop} again.

While these two programs produce the same result, they do it in very different ways. First, **fact_F** uses recursive types, whereas the **fact_T** does not. More importantly, **fact_F** uses a functional stack-based evaluation, whereas **fact_T** mutates registers and performs direct jumps. However, the proof of equivalence only differs in that we have to consider two cases - one in which they both diverge (for negative input *n*), and one in which they both terminate with related values (for non-negative input *n*).

6. Discussion and Future Work

FunTAL for Developers We have presented a multi-language **FT** that *safely* embeds assembly in a functional language. Moreover, our logical relation can be used to establish correctness of embedded assembly components. Developers of high-assurance software can write a high-level component **e** to serve as a specification for the TAL implementation **e** and use our logical relation to prove them equivalent.

FT also enables powerful compositional reasoning about high-level components, even in the presence of embedded assembly code. In fact, we conjecture that if the programmer does not use of stack-modifying lambdas, and if the embedded TAL contains no *statically defined* mutable tuples, then **FT** ensures referential transparency for high-level terms. Intuitively, in the absence of these side-channels (stack-manipulation and mutable cells), there is no way for two embedded TAL components to communicate with one another. Thus, even if a high-level term **e** contains embedded assembly, evaluating **e** has no observable effects. If the programmer does use stack-modifying lambdas or statically defined mutable tuples, reasoning about high-level components remains similar to reasoning about components in ML.

Continuation-Passing F and Rust Instead of trying to bridge the gap between the direct-style **F** and the continuation-aware **T**, we could have made **F** a continuation-passing-style

language, effectively lowering its level of abstraction to simplify interoperability with assembly. But the resulting multi-language would be more difficult for source programmers to use, as it would require them to reason about CPS'd programs. This is essentially the approach taken by the RustBelt project [?]—i.e., working with a CPS-style Rust with embedded unsafe C.¹ The project seeks to establish soundness of Rust and its standard library, where the latter essentially contains unsafe embedded C. In contrast to **T**, RustBelt does not take a multi-language approach or aim to handle inline assembly. Rather, it uses a sophisticated program logic for mutable state to reason about unsafe C code. It would be interesting to investigate a multi-language system with direct-style Rust interoperating with unsafe C and assembly along the lines of our work.

Choices in Multi-Language Design There are many potential choices when designing a multi-language system. For instance, we chose to expose low-level abstractions to high-level code by adding stack-modifying lambdas to **FT**, enabling more interactions between **F** and **T** code by invalidating equivalences that might otherwise have been used to justify correctness of compiler optimizations. We could also add foreign pointers to **FT**, which would allow references to mutable **T** tuples to flow into **F** as opaque values of lump type (as in Matthews-Findler [?]), allowing them to be passed but only used in **T**. Foreign pointers would have the form $\text{L}(\overline{\tau})\mathcal{FT}\ell$ (where $\ell : \text{ref } \langle \overline{\tau} \rangle \mathcal{T}$). While we can provide limited mutation to **F** via **T** libraries, foreign pointers would make that more flexible, albeit at the cost of complicating the multi-language.

Compositional Compiler Correctness As mentioned in §1, Perconti and Ahmed [?] proved correctness of a functional-language compiler that performs closure conversion and heap allocation. We can easily adapt our multi-language to verify correctness of a code-generation pass from their allocation target **A** to **T**, changing **FT** to **AT**. The semantics of **T** and **T**-relevant proofs in the logical relation can be reused without change. Correctness of code generation would then be expressed as contextual equivalence (\approx^{ctx}) in **AT**: if $e_A : \tau_A$ compiles to e_T then $e_A \approx^{ctx} \tau_A \mathcal{FT} e_T$.

JIT Formalization We plan to investigate modeling a JIT compiler using multi-language programs. The high-level source language would be untyped and the low-level language would be typed assembly (since type information is precisely what a JIT runtime discovers about portions of high-level code, triggering compilation). We would consider the space of JIT optimization to be the set of possible replacements of untyped components with sound low-level versions, with appropriate guards included to handle violation of typing assumptions. Note, of course, that the low-level versions may still have calls back into high-level untyped code. What the JIT is then doing at runtime is moving

¹ Personal communication with Derek Dreyer and Ralf Jung.

between those configurations, usually by learning enough type information to make the guards likely to pass.

We can prove a JIT compiler correct based on the transformations that it would do. Formally, for all moves between configurations that the JIT may perform, we must show:

$$\forall E, \mathbf{e}_S. \mathbf{e}_S \xrightarrow{E} \mathbf{e}_T \implies E[\mathbf{e}] \approx E[\mathcal{FT} \mathbf{e}_T]$$

where \xrightarrow{E} represents context-aware JIT-compilation that allows the compiler to use information in the context E , which could include values in scope, etc, in order to decide how to transform a component \mathbf{e}_S into \mathbf{e}_T . The definition of the JIT is thus \xrightarrow{E} , and we would prove equivalence of the resulting multi-language programs using a logical relation similar to the one shown in this paper.

7. Related Work

There has been a great deal of work on multi-language systems, typed assembly languages, logics for modular verification of assembly code, and logical relations in general. We focus our discussion on the most closely related work.

Our work builds on results about typed assembly [?] and in particular STAL, its stack-based variant [?]. §3 explains in detail the differences between our TAL and STAL. Note here though, that these differences stem from our goal to use type structure to define the notion of a TAL component. We share this goal with a number of previous type system design and verification efforts for flavors of assembly-like languages. [?] tackle the problem of safe linking for TAL program fragments and provide an extension of TAL’s type system that guarantees linking preserves type safety. [?] introduces a typed Floyd-Hoare logic for a stack-based low-level language that treats program fragments and their linking in a modular fashion. Outside the distinct technical details of what a component is in our TAL, our work differs from these results in that our notion of a TAL component matches that of a function in a high-level functional language.

Our multi-language semantics builds of work by Matthews and Findler [?] who gave multiple interoperability semantics between a dynamically and statically typed language. We also build on multi-languages used for compiler correctness [?] which embed the source (higher-level) and target (lower-level) languages of a compiler, though none of that work considers interoperability with a language as low-level as assembly.

A related strand of research explores type safety and foreign function interfaces (FFI). [?] describe sound type inference for the OCaml/C and JNI FFIs. [?] use a mixture of dynamic and static checks to construct a type-safe variant of JNI. [?] aim for fully abstract and type-safe interoperability between ML and a low-level language. However, their model low-level language is Scheme with reflection. [?] describes a core model for JNI that mixes Java bytecode and assembly. As an application, they design a sound type system for

their multi-language. Our work is distinct as it captures how assembly interacts safely with a functional language.

Our logical relation resembles the multi-language relation of Perconti-Ahmed [?] though theirs, without assembly, is simpler. Most prior logical relations pertaining to assembly or SECD machines are cross-language relations that specify equivalence of high-level (source) code and low-level (target) code and are used to prove compiler correctness [? ? ?]. Hur and Dreyer use a cross-language Kripke logical relation between ML and assembly to verify a one-pass compiler [?]. Neis *et al.* set up a parametric inter-language simulation (PILS) relating a functional source S and a CPS-style intermediate language I , and one relating I to a target assembly T [?]. None of these can reason about equivalence of (multi-block) assembly components as we do. Jaber and Tabareau [?] present a logical relation indexed by source-language types but inhabited by SECD terms, capturing high-level structure. Besides being able to reason about mixed programs, our \mathcal{FT} logical relation—indexed by multi-language types—is more expressive: it can be used to prove equivalence of assembly components of type τ when $\tau = \tau^T$ for some τ (analogous to Jaber-Tabareau) as well as when τ is *not* of translation type. All of these logical relations make use of biorthogonality, a natural choice for continuation-based languages.